

# Knowledge engineering

Basic knowledge modeling and RDF graph storage

## Semantic Web Link Open Data



Dr. Azanzi Jiomekong

University of Yaounde I, Department  
of Computer Science

10 avril 2023

Copyright (c) 2020 Azanzi Jiomekong.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.2 or any later version published by the Free Software Foundation ; with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts. More information about the license is included in the section entitled " The GNU Free Documentation License" .

# Questions

---

- How to represent simple knowledge on the Web ?
- How can we access the data we have already defined ?
- How to store RDF data ?

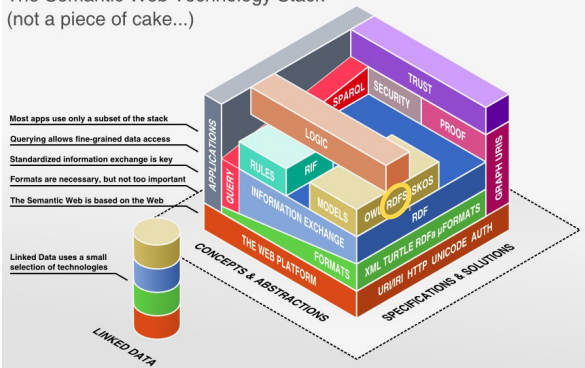
**The responses to these questions are given in this course**





# RDF Schema Language

The Semantic Web Technology Stack  
(not a piece of cake...)





# RDFSchema

---

- First W3C draft in April 1998,  
W3C Recommendation Feb. 2004.
- Defines a **data model** for the creation of RDF statements
- Allows :
  - Definition of **classes**  
**Class instantiation** in RDF via `< rdf : type >`
  - Definition of **properties** and **restriction**
  - Definition of **hierarchies**  
Subclasses and Superclasses  
Subproperties and Superproperties





# RDF Schema Language

## Vocabulary

---

- **rdfs :Class**

Concepts of a class, defines an abstract object and is applied (with `rdf :type`) to create instances

- **rdf :Property**

Base class for properties

- **rdfs :Literal**

Class for literals

## RDF Schema Language

## Vocabulary

- **rdfs :Resource**  
every entity of an RDF model is instance of this class
- Additionally  
**rdfs :Datatype**, **rdf :XMLLiteral**, **rdfs :Container**,  
**rdfs :ContainerMembershipProperty**

## Vocabulary



# RDF Schema Language

## Vocabulary : Properties

---

- **rdfs :subClassOf**

Transitive property to define inheritance hierarchies for classes

- **rdfs :subProperty**

Transitive property to define inheritance hierarchies for properties

- **rdfs :domain**

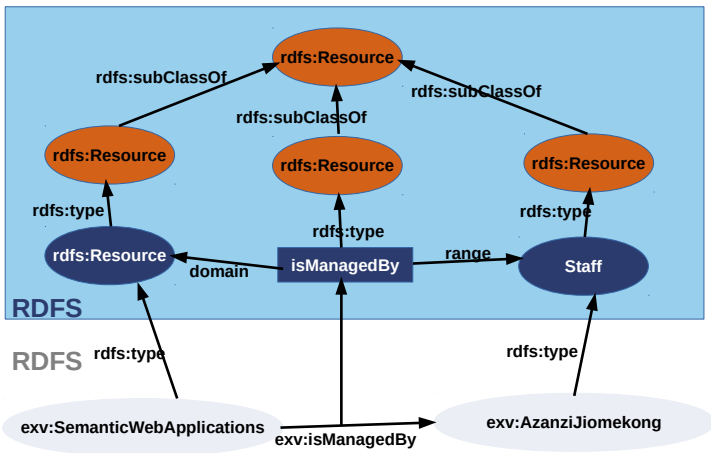
defines the domain of a property concerning a class

- **rdfs :range**

Defines range of a property concerning a class

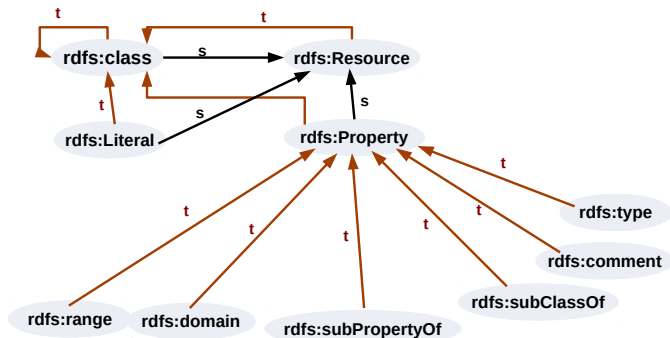
# RDF Schema Language

## Vocabulary



# RDF Schema Language

## Language Model



s: subclass relation

t: instance relation

# RDF Schema Language

Vocabulary : further properties

---

- **rdfs :seeAlso**

Defines a relation of a resource to another, which explains it

- **rdfs :isDefinedBy**

subproperty of **rdfs :seeAlso**, defines the relation of a resource to its definition

- **rdfs :comment**

Comment, usually as text

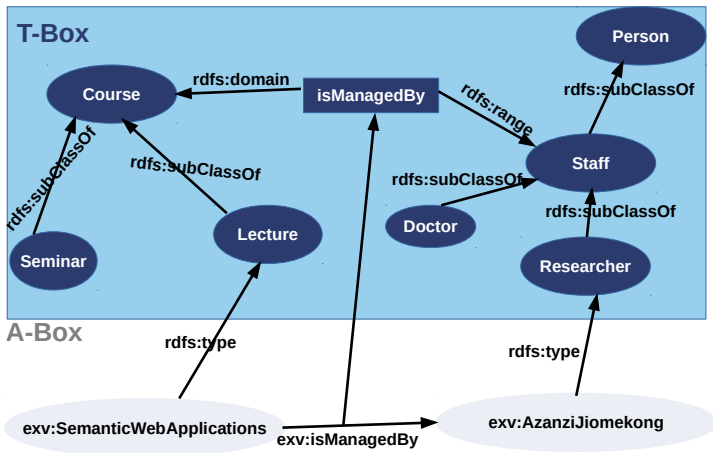
- **rdfs :label**

Readable name of a resource (contrary to ID)



# RDF Schema Language

## Knowledge base example



# RDF Schema Language

## Knowledge base example

---

```
@prefix rdf:<http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
@prefix rdfs:<http://www.w3.org/2000/01/rdf-schema#> .

:Course a rdfs:Class .
:Lecture a rdfs:Class ;
    rdfs:subClassOf :Course .
:Seminar a rdfs:Class ;
    rdfs:subClassOf :Course .

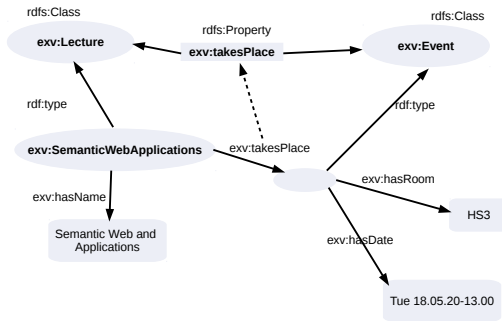
:Person a rdfs:Class .
:Staff a rdfs:Class ;
    rdfs:subClassOf :Person .
:Researcher a rdfs:Class ;
    rdfs:subClassOf :Staff .
:Doctor a rdfs:Class ;
    rdfs:subClassOf :Staff .

:isManagedBy a rdf:Property ;
    rdfs:domain :Course ;
    rdfs:range :Staff .

:SemanticWebApplications a :Lecture .
:AzanziJiomakong a :Researcher .
:SemanticWebApplications :isManagedBy :AzanziJiomakong .
```

# Semantics is the in RDF and RDFS ?


Let's consider the following knowledge base



The semantics of a term from an RDF(S) ontology is given in terms of its properties and its values (instances)

# Semantics is the in RDF and RDFS


What conclusions can we deduce with RDF(S) ?



- Deduction of entity class membership from the domain of one of its properties
- We can deduce that Semantic Web Technologies must be a lecture

# Semantics is the in RDF and RDFS

What conclusions can we deduce with RDF(S) ?

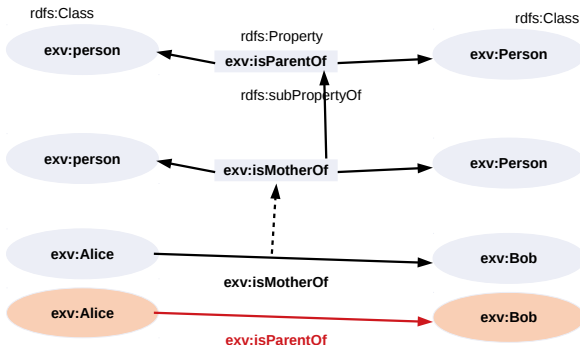


- Deduction of entity class membership from the range of its properties
- We can deduce that the blank node must represent an event



# Semantics is the in RDF and RDFS

What conclusions can we deduce with RDF(S) ?



Deduction of new facts from subproperty relationships

# RDF Schema Language

---

- Small vocabulary
- We don't have something like a negation
- It is rather difficult to create inconsistencies and failure
- Everything define with RDFS is potentially true
- What is missing is the formal vocabulary
- For machine to understand, we need a formal representation with logic



# Question

---

How to store RDF triples to carry out efficient SPARQL queries?



# Storing RDF data

## RDBMS

---

- RDF graphs can be stored in a set of triples (s, p, o)
- Triples can be stored in a RDBMS
- Three steps for SPARQL query processing :
  - Convert SPARQL query to SQL query
  - Use RDBMS to answer SQL query
  - Generate SPARQL query result from SQL query result



# Storing RDF data

## Giant Triple Storage

---

- Basic idea
  - Store all RDF triples in a single table
  - Translate SPARQL query into SQL query
  - Performance depends on efficient indexing
- Advantages : easy to implement, works for huge numbers of properties is indexes are chosen with care
- Disadvantage : many self joins

# Storing RDF data

## Giant Triple Storage

### Subject Property Object

|     |            |               |
|-----|------------|---------------|
| ID1 | type       | FullProfessor |
| ID1 | teacherOf  | INFO201       |
| ID1 | bachelorFr | UDS           |
| ID1 | mastyerFr  | UY1           |
| ID1 | phdFr      | UY1           |
| ID2 | type       | AssocProf     |
| ID2 | worksFor   | UY1           |
| ID2 | teacherOf  | INF4288       |
| ID2 | bachelorFr | UMA           |
| ID2 | phdFr      | UY1           |
| ID3 | type       | GradStudent   |
| ID3 | advisor    | ID2           |
| ID3 | teachAss   | INFO201       |
| ID3 | bachelorFr | UY1           |
| ID3 | mastyerFr  | UY1           |
| ID4 | type       | GradStudent   |
| ID4 | advisor    | ID1           |
| ID4 | TakesCour  | SWA           |
| ID4 | bachelorFr | UMA           |

# Storing RDF data

## ID Based Triple Storage

- Use numerical identifier for each RDF term in the dataset
- Saves space and enhances efficiency

| RDF Term       | ID | s  | p  | o  |
|----------------|----|----|----|----|
| :ID1           | 1  | 1  | 2  | 3  |
| rdf:type       | 2  | 1  | 4  | 5  |
| :FullProfessor | 3  | 1  | 6  | 7  |
| :teacherOf     | 4  | 1  | 8  | 9  |
| "SWA"          | 5  | 1  | 10 | 11 |
| :bachelorFrom  | 6  | 12 | 13 | 14 |
| "UDS"          | 7  | 12 | 15 | 7  |
| :masterFrom    | 8  | 12 | 4  | 16 |
| "UY1"          | 9  |    |    |    |





# Storing RDF data

## Properties Tables

- Combining all (or some) properties of similar subjects in n-ary tables
- Use ID based encoding for efficiency

### Professors

| ID   | type           | teacherOf | bachelorFrom | masterFrom | phdFrom | worksFor |
|------|----------------|-----------|--------------|------------|---------|----------|
| ID 1 | FullProfessor  | "Algo"    | "UY1"        | "UY1"      | "UMA"   | NULL     |
| ID 2 | AssocProfessor | "OOP"     | "UMA"        | NULL       | "UDA"   | UY1      |

### Students

| ID   | type        | advisor | bachelorFrom | masterFrom | teachingAss | takesCourses |
|------|-------------|---------|--------------|------------|-------------|--------------|
| ID 3 | GradStudent | ID2     | UDA          | UY1        | SWA         | NULL         |
| ID 4 | GradStudent | ID1     | UDA          | NULL       | "NULL"      | "OOP"        |

# Storing RDF data

## Properties Tables

---

### Advantages :

- Fewer joins
- if the data is structures, we have a relational DB

### Disadvantages

- Potentially a lot of NULLs
- Clustering is not trivial
- Multi-value properties are complicated

# Storing RDF data

## Vertically Partitioned Tables (Binary Tables)

- For each unique property create a two column table
- Use ID based encoding for efficiency

| type |                | teacherOf | bachelorFrom |       |
|------|----------------|-----------|--------------|-------|
| ID 1 | FullProfessor  |           | ID 1         | "UY1" |
| ID 2 | AssocProfessor | "Algo"    | ID 2         | "UMA" |
| ID 3 | GradStudent    | "OOP"     | ID 3         | "UY1" |
| ID 4 | GradStudent    |           | ID 4         | "UDA" |

| masterFrom |       | phdFrom | worksFor |
|------------|-------|---------|----------|
| ID 1       | "UY1" | ID 1    | "UMA"    |
| ID 3       | "UY1" | ID 2    | "UDA"    |

| advisor | bachelorFrom | teachingAss |
|---------|--------------|-------------|
| ID3 ID2 | ID 1 "UY1"   |             |
| ID4 ID1 | ID 2 "UMA"   | ID3 "SWA"   |
|         | ID 3 "UDA"   |             |
|         | ID 4 "UDA"   |             |

| takesCourses |
|--------------|
| ID4 "OOP"    |

# Storing RDF data

## Vertically Partitioned Tables (Binary Tables)

---

### Advantages :

- Support multi-value properties
- No NULLs
- Read only needed attributes (i.e. less I/O)
- No clustering
- Excellent performance (if number of properties is small, queries with bounded properties)

### Disadvantages :

- Expensive inserts
- Bad performance (large number of properties, queries with unbounded properties)


# Storing RDF data

## Vertically Partitioned Tables (Binary Tables)

- Different physical storage models for relational DBs
- Two types of storage : Row Based Storage and Column Based Storage

### Row Based Storage :

- Tuples (i.e. DB records) are stored consecutively
- Entire row needs to be read even if few attributes are projected




# Storing RDF data

## Vertically Partitioned Tables (Binary Tables)


### Column Based Storage :

- Read columns relevant to the query  
→ projection is free
- Inserts are expensive



## Storing RDF data

## Hexastores







# Storing RDF data

## Triple Stores

---

- AnzoGraph -  
<https://docs.cambridgesemantics.com/anzograph/v2.2/userdoc/home.htm>
- Apache Jena TDB -
- Open Link Virtuoso -
- AllegroGraph -
- GraphDB -
- Stardog -

W3C maintains a list of triple stores :

<https://www.w3.org/wiki/LargeTripleStores>

# Limits of RDF and RDFS

- RDF and RDFS are sufficient to represent the knowledge of the world ?
- What can be represented and what can't be represented ?
- How can we represent the following knowledge :
  - "Certain animals only eat meat. Animal like Cow only eats vegetables. Certain animals eat vegetables and meats" ? → Locality of global properties
  - "Man and human are subclasses of Human. A man is not a woman" → Disjunctive Classes
  - "Human has two parents" → Restrictions on properties - Cardinality restriction

# Limits of RDF and RDFS

## Locality of global properties



## Problem


- Cows only eat vegetables
- Other animals also eat meat



# Limits of RDF and RDFS

## Cardinality Restrictions

---



Problem : Every human has two parents

## Limits of RDF and RDFS

## Special Property Constraints

- **Transitivity** (e.g. "is greater than")
- **Uniqueness** (e.g. "is mother of")
- **Inversiveness** (e.g. "is parent of" and "is child of")



